

1 Capabilities as First-Class Modules with Separate Compilation

2 Jam Kabeer Ali Khan
3 jamkhan@connect.hku.hk
4 The University of Hong Kong
5 Hong Kong SAR, China

6 ABSTRACT

7 We present ENVCAP, a statically typed programming language
8 based on environment-based semantics that supports *first-class*
9 environments, capabilities, and separate compilation. By utilizing
10 the environment-based semantics of λ_E [15], ENVCAP models
11 capabilities [5] as *first-class* modules [15], as an alternative to
12 the object-capability model [12], and enables separate compilation
13 without extra-linguistic structures such as *linksets* [2].

14 1 PROBLEM & MOTIVATION

15 Programming language implementations often use environments—
16 maps of bindings and values—for efficiency, while theoretical calculi
17 rely on substitution, weakening the correctness guarantees of im-
18 plementations. Environment-based semantics [4, 8, 14, 15] address
19 this gap by incorporating environments into formal calculi, en-
20 abling *first-class* environments—environments as values that can be
21 manipulated. This approach provides a foundation for addressing
22 two key problems in language design and implementation: (1) ca-
23 pabilities are commonly modeled as objects [5], which introduces
24 complexity due to object-oriented calculi, making reasoning about
25 and implementing capabilities challenging; and (2) traditional ap-
26 proaches to separate compilation rely on external structures, such
27 as *linksets* [2], which are not part of the core language.

28 ENVCAP is the first programming language to address these
29 problems by providing two key solutions based on *first-class* envi-
30 ronments:

- 31 (1) **Capabilities as First-Class Modules:** ENVCAP models
32 capabilities as *first-class* modules [15] using *first-class* en-
33 vironments, offering a simpler alternative to the object-
34 capability model [12].
- 35 (2) **Separate Compilation:** ENVCAP enables separate com-
36 pilation entirely within the core λ_E language, eliminating
37 the need for extra-linguistic structures such as linksets.

38 2 BACKGROUND & RELATED WORK

39 The λ calculus relies on substitution-based semantics, where beta-
40 reduction replaces terms during application. This approach is ineffi-
41 cient in practice and challenging to formalize due to issues like name
42 capture. Environment-based semantics address these limitations by
43 integrating environments—maps of bindings and values—directly
44 into the formalization, aligning with implementation practices. The
45 λ_E [15] and E_i [14] calculi formalize this approach, unifying ex-
46 pressions and environments to enable *first-class* environments and
47 introducing the box construct $e_1 \triangleright e_2$, which evaluates e_2 under
48 environment e_1 , allowing for modeling capabilities as *first-class*
49 modules.

50 Capabilities [5] enforce access control, commonly implemented
51 using the object-capability model in programming languages such

52 as Newspeak [1] and Wyvern [7, 11]. As their name implies, ca-
53 pabilities are primarily designed to restrict access to resources.
54 When modeled as first-class modules, capabilities can be passed as
55 arguments, effectively granting access to specific resources. Tan
56 and Oliveira [15] proposed sandboxed *first-class* modules as an
57 alternative, where sandboxed modules restrict access to the global
58 environment by requiring all dependencies to be explicitly passed
59 as arguments, allowing to model capabilities. In addition, most
60 capability-based languages lack formalization for separate compila-
61 tion, which is essential to maintain authority control.

62 Separate compilation enables programs to type-check and com-
63 pile based on interfaces rather than implementations. Cardelli [2]
64 introduced *linksets* as a formal framework for separate compilation,
65 later adapted by Standard ML [13]. In this framework, *linksets*
66 provides a structure to link code components post-compilation, for-
67 malized through a simple module system called *bindings*. However,
68 *linksets* are extra-linguistic—they are not part of the core language
69 and lack concrete syntax, complicating implementation.

70 A simpler and more generalized solution could leverage core lan-
71 guage constructs, where program fragments act as abstractions over
72 interfaces, and interfaces are unified with types [14]. For instance,
73 if component A is imported by component B , it can be represented
74 as $\lambda(\text{interface } A).B$. Using *first-class* environments in λ_E , ENVCAP
75 enables separate compilation within the core language, eliminating
76 the need for extra-linguistic mechanisms. Additionally, Cardelli’s
77 framework lacks support for *first-class* modules. This work extends
78 Cardelli’s framework by replacing *linksets* with constructs from the
79 λ_E calculus, offering a more expressive and streamlined solution.

80 3 APPROACH & UNIQUENESS

81 The semantics of ENVCAP are derived from λ_E via elaboration and
82 syntax design is inspired by OCaml [9], Wyvern [7, 11], and λ_E .

Fragment	$U ::= \text{program } \langle S, I, R, E, A \rangle$	100
Authority	$S ::= @\text{pure} \mid @\text{resource}$	101
Import	$I ::= \cdot \mid \text{import } l : A, I$	102
Requirements	$R ::= \cdot \mid \text{require } l : A, R$	103
Expressions	$E ::= \text{env} \mid E.n \mid i \mid e \mid \lambda A. E \mid \text{with } E_1 \text{ in } E_2$ $\mid E_1 E_2 \mid E_1; E_2 \mid E_1, E_2 \mid \{\ell = E\} \mid E.\ell$ $\mid \text{function } \ell A : B E \mid \text{struct } A E \mid \text{struct } E$ $\mid E_1 * E_2 \mid \text{functor } \ell A : B E \mid \text{module } \ell : B E$ $\mid \text{let } x E_1 \mid \text{open } E_1 E_2 \mid E : A$	104
Types	$A, B, \Gamma ::= \text{Int} \mid e \mid A \rightarrow B \mid \{\ell : A\}$ $\mid A \& B \mid \text{Sig}[A, B]$	105

114 **Figure 1: Core ENVCAP syntax.**

3.1 Syntax & Design

Figure 1 presents the core syntactic constructs of ENVCAP and we discuss the most relevant constructs. A fragment (program $S \ I \ R \ E \ A$) represents an implementation paired with an interface A , where interfaces desugar into types. First-class modules include struct $A \ E$ (taking A as input) and struct E (syntactic sugar for E). The construct module $\ell : B \ E$ desugars into a record $\{\ell = E : B\}$, while functor $\ell \ A : B \ E$ desugars into $\{\ell = \text{struct } A \ (E : B)\}$. Module application $(E_1 * E_2)$ requires E_1 to have type $\text{Sig}[A, B]$. The open $E_1 \ E_2$ construct requires E_1 to be of record type and loads its contents into the current context via elaboration. The reification operator env retrieves the current environment. Sequences $E_1; E_2$ elaborate to dependent merges (e_1, e_2) in λ_E , which is helpful to model a sequence of declarations, e.g. $\{x = 1\}, \{y = x + x\}$, allowing E_2 to depend on E_1 . Pairs (E_1, E_2) are non-dependent and generalized to support *tuples*, but elaborate to e_1, e_2 in λ_E . The construct function $\ell \ A : B \ E$ elaborates into fix $A \rightarrow B.E$, enabling recursion.

3.2 First-class Environments & Modules

We illustrate *first-class* environments and modules with an example. First, we define a @pure fragment with *UTIL* and *MATH* interfaces:

```
1 @pure module Example1
2 interface UTIL {val diff : Int;};
3 interface MATH {val fact : Int -> Int};
```

Next, we define a functor (a parameterized module) that desugars into a labeled record with an anonymous *first-class* module (*struct A E*):

```
1 functor math (util: UTIL) : MATH {
2   (* loads contents into environment, e.g. diff *)
3   open util;
4   function fact(n: Int): Int {
5     if (n == 0) then 1 else n * fact(n - diff)
6   }
7};
```

Finally, we use the *with E₁ E₂* construct to evaluate E_2 in environment E_1 . We create a new environment containing the current environment (extracted via the *env* operator), the result of applying the *math* functor to an anonymous module of type *UTIL*, and a new binding for *x*. The computation proceeds under this new environment, computing factorials for both the old and new values of *x*:

```
1 let x = 5;
2 with
3 ({prevEnv = env}; math(struct {let diff = 1}); {x = 6})
4 in { let resultOld = fact(prevEnv.x); (* 120 *)
5   let resultNew = fact(x)           (* 720 *)}
```

This example demonstrates ENVCAP’s support for *first-class* modules and environments, showcasing its expressive power.

3.3 Capabilities as First-Class Modules

Capabilities in ENVCAP are modeled using **sandboxed first-class modules**, which enforce controlled access by requiring all resources to be explicitly passed as parameters, and **authority annotations** (@pure and @resource) inspired by Wyvern [11]. The modules are sandboxed via elaboration to $\epsilon \triangleright \lambda A.e$ in λ_E , so computation runs under *empty* (ϵ) environment and hence, resources

are only passed via parameters. A @resource fragment can import any fragment, while a @pure fragment can only import other @pure fragments. If a @pure fragment requires functionality from a @resource fragment, it must declare it as a requirement rather than an import, ensuring explicit capability propagation. ENVCAP supports the use of interface files in requirements, so interfaces of required fragments can be utilized for separate compilation. Only valid imports are linked and the validity of imports with authority annotations is checked during elaboration to λ_E , maintaining authority control.

We present an example of how capabilities are propagated in ENVCAP, consisting of a pure fragment B and a resource fragment A.

```
1 @pure module B
2 require (U: System.Utils);
3 let mapList =
4   U.Map(\(x:Int) => {x + 1}, [1, 2, 3])
```

Fragment B is a @pure module requiring *System.Utils* as a capability. Since *System.Utils* is a resource, it cannot be imported directly. Instead, it is passed as a parameter, ensuring explicit capability propagation.

```
1 @resource module A
2 import System.Utils B;
3
4 let result = B(System.Utils).mapList
```

Fragment A, a @resource module, imports *System.Utils* and instantiates B with it. The *mapList* function is extracted and assigned to the result. Here, *System.Utils* acts as a capability: A has the authority to access and pass it to B, which requires this capability.

3.4 Separate Compilation with Unified Types and Interfaces

We enable separate compilation by unifying interfaces and types, where an implementation file is treated as an abstraction over the interfaces of its imports at the core level of λ_E . The linking mechanism is also defined at the core level, eliminating the need for external structures.

3.4.1 Example. To illustrate this, we adapt one of Cardelli’s examples [2] in ENVCAP. For clarity, we present an example within a single file.

```
1 @pure module Example3
2 interface N { val x : Int };
3 module n : N {
4   let x = 3
5 }
```

n is a module that implements interface N. Next, we define a functor m that implements interface M and takes an implementation of N as input; this is similar to an import at the fragment level.

```
1 interface M {
2   val f : Int -> Int;
3   val m : Int
4 };
5 functor m (n: N) : M {
6   open n;
7   let f = \ (y: Int) => y + x;
8   let m = f(x)
```

175 9 }
 176 In this example, the functor m represents a program fragment
 177 that imports the module n with the interface N . In ENVCAP, inter-
 178 faces are treated as types, so *interface N* and *interface M* desugar
 179 into record types $\{N : \{x : \text{Int}\}\}$ and $\{M : \{f : \text{Int} \rightarrow \text{Int}\} \& \{m : \text{Int}\}\}$, respectively.
 180 Int }
 181

182 3.4.2 *Compilation: ENVCAP $\rightsquigarrow \lambda_E$* . In our setting, the elaboration
 183 from ENVCAP to λ_E is analogous to Cardelli's compilation of bind-
 184 ings to linksets. Modules and functors elaborate into records and
 185 boxed abstractions. In this example, specifically:

```
186    module n : {n : {x : Int}} ~> {n = {x = 3} : {x : Int}}  

  187  

  188    functor m : {m : sig[ {n : {x : Int}}, {f : Int → Int} & {m : Int} ]}  

  189  

  190    ~> {m = ε > λ{n : {x : Int}}.(f = .., m = (?f)(?.n.x)) : type..}
```

191 For simplicity, we omit the elaboration of the open (given in
 192 Figure 2) statement, which essentially loads the contents of module
 193 n ($\{x = 3\}$) into the environment.

194 3.4.3 *Linking*. Once elaborated to λ_E , the expressions can be linked.
 195 To ensure correct linking order, we use Kahn's [6] algorithm for
 196 topological sorting to determine module dependencies and linking
 197 order. It enables the detection of cyclic dependencies that are not
 198 allowed in ENVCAP for both imports and requirements for the sake
 199 of simplicity. The linking process proceeds as follows.

- 200 (1) **Initial State:** $L \equiv \epsilon$.
- 201 (2) **Link n:** $L \hookleftarrow L' \equiv \epsilon, \{n = {x = 3} : {x : \text{Int}}\}$.
- 202 (3) **Link m:** $L' \hookleftarrow L', \{m = \epsilon, \{n = {x = 3}\} \triangleright (f = .., m =$
 203 $(?.f)(?.n.x)) : \{f : \text{Int} \rightarrow \text{Int}\} \& \{m : \text{Int}\}\}$.

204 When no further linking steps are possible ($L \not\hookleftarrow$), the resulting
 205 structure, with dependent merges (\triangleright), contains fully linked frag-
 206 ments that can be executed individually.

4 RESULTS & CONTRIBUTIONS

211 The main results achieved are the following: 1) Implementation
 212 of the ENVCAP interpreter in Haskell [10]; 2) Formalization of
 213 type-directed elaboration $\text{ENVCAP} \rightsquigarrow \lambda_E$ in Rocq [16].

4.1 Implementation

214 The ENVCAP interpreter, implemented in ~ 5000 lines of Haskell,
 215 supports recursion, algebraic datatypes, and other key features. Its
 216 architecture comprises parsing, locally nameless transformation,
 217 desugaring, elaboration, and execution. The locally nameless repre-
 218 sentation [3] simplifies implementation by avoiding ambiguous en-
 219 vironment lookups, which λ_E forbids. Furthermore, the interpreter
 220 produces .epc files containing λ_E expressions, enabling indepen-
 221 dent linking and execution, ensuring flexibility by decoupling from
 222 specific code generation.

4.2 Meta-theory

223 Type-directed elaboration translates *core* ENVCAP to λ_E expres-
 224 sions while preserving type consistency. Currently, the proof ex-
 225 cludes the compilation of fragments of the form program $\langle S, I, R, E, A \rangle$;
 226 the key elaboration rules are shown in Figure 2.

EL-MODULE	$\epsilon \& A \vdash E : B \rightsquigarrow e$
	$\Gamma \vdash \text{struct } A E : \text{Sig}[A, B] \rightsquigarrow \epsilon \triangleright \lambda A .e$
EL-MODAPP	$\Gamma \vdash E_1 : \text{Sig}[A, B] \rightsquigarrow e_1 \quad \Gamma \vdash E_2 : A \rightsquigarrow e_2$
	$\Gamma \vdash E_1 * E_2 : B \rightsquigarrow e_1 e_2$
EL-CLOS	$\Gamma \vdash E_1 : \Gamma_1 \rightsquigarrow e_1 \quad \Gamma_1 \& A \vdash E_2 : B \rightsquigarrow e_2$
	$\Gamma \vdash \langle E_1, \lambda A. E_2 \rangle : A \rightarrow B \rightsquigarrow e_1 \triangleright \lambda A .e_2$
EL-NON-DEPENDENT-MERGE	$\Gamma \vdash E_1 : A_1 \rightsquigarrow e_1 \quad \Gamma \vdash E_2 : A_2 \rightsquigarrow e_2$
	$\Gamma \vdash E_1, E_2 : A_1 \& A_2 \rightsquigarrow (\lambda \Gamma .(\underline{0} \triangleright e_1), (\underline{1} \triangleright e_2)) ?$
EL-OPEN	$\Gamma \vdash E_1 : \{\ell : A\} \rightsquigarrow e_1$
	$\Gamma \vdash E_1.l : A \rightsquigarrow e_1.l \quad \Gamma \& A \vdash E_2 : B \rightsquigarrow e_2$
	$\Gamma \vdash \text{open } E_1 E_2 : B \rightsquigarrow (\lambda A .e_2) (e_1.l)$

where $|.|$ is simply translation of ENVCAP types to λ_E types.

Figure 2: Elaboration: $\Gamma \vdash \text{ENVCAP} : A \rightsquigarrow \lambda_E$

The theorems formalized in Rocq are the following:

THEOREM 1 (TYPE PRESERVATION).

if $\Gamma \vdash E : A \rightsquigarrow e$, then $|\Gamma| \vdash e : |A|$.

THEOREM 2 (UNIQUENESS OF TYPE INFERENCE).

if $\Gamma \vdash E : A_1 \rightsquigarrow e_1$ and $\Gamma \vdash E : A_2 \rightsquigarrow e_2$, then $A_1 \equiv A_2$.

THEOREM 3 (UNIQUENESS OF ELABORATION).

if $\Gamma \vdash E : A_1 \rightsquigarrow e_1$ and $\Gamma \vdash E : A_2 \rightsquigarrow e_2$, then $e_1 \equiv e_2$.

Significance. ENVCAP presents *first-class* modules as an alterna-
 tive to objects and enables separate compilation without external
 constructs, such as linksets. This simplifies the implementation of
 programming languages by providing simpler alternatives using
first-class environments and environment-based semantics.

5 ONGOING AND FUTURE WORK

The ENVCAP project aims to formalize separate compilation in Rocq,
 based on Cardelli's framework. Inspired by WyVERN, it enforces
 authority safety for capability guarantees. To enhance modularity,
 it introduces subtyping support, enabling export restrictions via
 interface files.

ACKNOWLEDGEMENTS

We thank Bruno C. d. S. Oliveira and Jinhao Tan for their founda-
 tional work on first-class environments and supervision.

REFERENCES

- [1] Gilad Bracha, Peter von der Ahé, Vassili Bykov, Yaron Kashai, William Maddox, and Eliot Miranda. 2010. Modules as Objects in Newspeak. In *ECOOP 2010 - Object-Oriented Programming, 24th European Conference, Maribor, Slovenia, June 21-25, 2010. Proceedings (Lecture Notes in Computer Science, Vol. 6183)*, Theo D'Hondt (Ed.). Springer, 405–428. https://doi.org/10.1007/978-3-642-14107-2_20
- [2] Luca Cardelli. 1997. Program fragments, linking, and modularization. In *Proceedings of the 24th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (Paris, France) (POPL '97). Association for Computing Machinery, New York, NY, USA, 266–277. <https://doi.org/10.1145/263699.263735>
- [3] Arthur Charguéraud. 2012. The Locally Nameless Representation. *Journal of Automated Reasoning* 49, 3 (2012), 363–408. <https://doi.org/10.1007/s10817-011-9225-2>
- [4] Pierre-Louis Curien. 1991. An Abstract Framework for Environment Machines. *Theor. Comput. Sci.* 82 (1991), 389–402. <https://api.semanticscholar.org/CorpusID:41601809>
- [5] Jack B. Dennis and Earl C. Van Horn. 1966. Programming semantics for multiprogrammed computations. *Commun. ACM* 9, 3 (March 1966), 143–155. <https://doi.org/10.1145/365230.365252>
- [6] A. B. Kahn. 1962. Topological sorting of large networks. *Commun. ACM* 5, 11 (Nov. 1962), 558–562. <https://doi.org/10.1145/368996.369025>
- [7] Darya Kurilova, Alex Potanin, and Jonathan Aldrich. 2014. Wyvern: Impacting Software Security via Programming Language Design. In *Proceedings of the 5th Workshop on Evaluation and Usability of Programming Languages and Tools* (Portland, Oregon, USA) (PLATEAU '14). Association for Computing Machinery, New York, NY, USA, 57–58. <https://doi.org/10.1145/2688204.2688216>
- [8] Peter J. Landin. 1964. The Mechanical Evaluation of Expressions. *Comput. J.* 6 (1964), 308–320. <https://api.semanticscholar.org/CorpusID:5845983>
- [9] Xavier Leroy, Damien Doligez, Alain Frisch, Jacques Garrigue, Didier Rémy, and Jérôme Vuillon. 2020. *The OCaml system release 4.10*. <https://caml.inria.fr/pub/docs/manual-ocaml/>
- [10] Simon Marlow (Ed.). 2010. *Haskell 2010 – Language Report*. www.haskell.org/onlinereport/haskell2010/
- [11] Darya Melicher, Yangqingwei Shi, Alex Potanin, and Jonathan Aldrich. 2017. A Capability-Based Module System for Authority Control. In *31st European Conference on Object-Oriented Programming (ECOOP 2017) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 74)*, Peter Müller (Ed.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 20:1–20:27. <https://doi.org/10.4230/LIPIcs.ECOOP.2017.20>
- [12] Mark Samuel Miller. 2006. *Robust Composition: Towards a Unified Approach to Access Control and Concurrency Control*. Ph.D. Dissertation. Johns Hopkins University, Baltimore, Maryland, USA.
- [13] David Swasey, Tom Murphy, Karl Crary, and Robert Harper. 2006. A separate compilation extension to standard ML. In *Proceedings of the 2006 Workshop on ML* (Portland, Oregon, USA) (ML '06). Association for Computing Machinery, New York, NY, USA, 32–42. <https://doi.org/10.1145/1159876.1159883>
- [14] Jinhao Tan and Bruno C. d. S. Oliveira. 2023. Dependent Merges and First-Class Environments. In *37th European Conference on Object-Oriented Programming (ECOOP 2023) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 263)*, Karim Ali and Guido Salvaneschi (Eds.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 34:1–34:32. <https://doi.org/10.4230/LIPIcs.ECOOP.2023.34>
- [15] Jinhao Tan and Bruno C. d. S. Oliveira. 2024. A Case for First-Class Environments. *Proc. ACM Program. Lang.* 8, OOPSLA2, Article 360 (Oct. 2024), 30 pages. <https://doi.org/10.1145/3689800>
- [16] The Coq Development Team. 2024. The Coq Reference Manual – Release 8.19.0. <https://coq.inria.fr/doc/V8.19.0/refman.html>